

ADA IA 2

Netflix Personalized Movie and Show Recommendation

Abstract

Netflix is a strong real-world example of how data analytics supports personalization in a large digital platform. Its challenge is not merely to store a large catalogue of films and shows, but to help each member discover relevant content quickly. This report studies Netflix's recommendation problem and maps it primarily to **Collaborative Filtering**, including user-based filtering, item-based filtering and Matrix Factorization.

Collaborative Filtering remains an important conceptual and historical foundation because it learns relationships between member profiles and titles from interaction data. Nevertheless, it is not sufficient by itself for Netflix's current scale. Netflix's recent technical work shows an evolution toward transformer-based recommendation foundation models, generative recommendation and LLM-based ranking. These modern systems can represent long-term histories, recent activity, title metadata and contextual signals more effectively than a static user-item matrix.

The report critically evaluates the advantages, limitations, computational complexity, scalability, cold-start problem, popularity bias, diversity and evaluation of Netflix's recommendation approach. It concludes that the most suitable architecture is hybrid: Collaborative Filtering can generate useful candidates, content-based methods can help with new titles, and sequential or foundation models can perform contextual ranking and homepage construction.

1. Company Overview

1.1 Brief introduction

Netflix is a global entertainment technology company offering films, television series, games, live programming and other digital entertainment. It began in 1997 as a DVD-rental company and moved into streaming in 2007.

The platform's main personalization challenge is content discovery. A member may have access to thousands of titles, but only a limited number can be displayed prominently on the screen. Netflix must therefore decide which titles to show, how to rank them and how to organize them into rows.

Netflix's recommendation technology can personalize:

- Titles displayed on the homepage.
- The order of titles within a row.
- Rows such as "Because You Watched."
- "Continue Watching" recommendations.
- Search and browsing results.
- Artwork and visual presentation.
- The balance between familiar titles and new discoveries.

Netflix's official recommendation-help page explains that its systems rank and position titles on the homepage using multiple layers of personalization.




Figure 1. Example of a personalized Netflix homepage.

1.2 Scale of operations

Netflix's scale makes recommendation a major data-analytics and systems-engineering problem.

Netflix reported approximately **\$45.18 billion in revenue for 2025**, representing 16% year-over-year growth. It also reported more than **325 million paid memberships at the end of 2025**. However, membership totals should be used carefully because Netflix changed its reporting approach and does not provide membership figures with the same frequency as in earlier years.[\[netflixtechblog\]](#)[\[netflixtechblog\]](#)

A membership may also represent multiple people in a household, while different profiles may have different tastes. The recommendation system must therefore personalize at the profile level rather than assume that every person using one account has identical preferences.

At this scale, Netflix's recommendation system must:

- Process very large volumes of behavioural data.
- Support members in different countries and languages.
- Respect different catalogues by region.
- Respond with low latency.
- Adapt to changing interests.
- Support multiple devices.
- Balance relevance with diversity.
- Handle new members and newly released titles.

Table 1. Why Netflix is a large-scale analytics problem

Challenge	Effect on recommendation design
Hundreds of millions of memberships	Models must operate across a very large user population
Large and changing catalogue	Recommendations must be refreshed frequently
Sparse viewing history	Members interact with only a small fraction of the catalogue
Multiple profiles per household	Recommendations must be profile-specific
Global availability	Language, region and cultural context affect relevance
Real-time browsing	Results must be generated quickly
New titles and new members	Cold-start strategies are required

1.3 Data generated by Netflix

Netflix's recommendation models learn from interactions between members and the platform.

Behavioural information

Examples include:

- Titles watched.
- Watch duration.
- Completion or abandonment.
- Searches.
- Browsing and scrolling.
- Thumbs-up or other feedback.
- Repeated viewing.
- Recent activity.
- "Continue Watching" behaviour.
- Add-to-list activity.

Content information

Examples include:

- Title.
- Movie or series classification.
- Genre.
- Language.
- Country of production.
- Release year.
- Story and thematic information.
- Cast and creators.
- Artwork.

Contextual information

Examples include:

- Member profile.
- Device type.
- Time of day.
- Locale.
- Language.

- Current page or surface.
- Recent browsing activity.

Netflix's Foundation Model publication describes the use of long-term interaction histories, content data and request-time context in a shared recommendation model.[\[netflixtechblog\]](#)

It is important to distinguish between raw events, processed interactions and model tokens. Netflix's public sources do not establish a current official value for "petabytes of recommendation data generated per day." Unsupported third-party estimates should therefore not be included.

2. Business Problem

2.1 What problem is being solved?

Netflix is solving a **personalized content-discovery and ranking problem**.

The system must answer questions such as:

- Which titles are most relevant to this member?
- Which recommendation rows should be displayed?
- Which titles should appear in each row?
- In what order should titles be shown?
- How should recent activity affect recommendations?
- Which watched titles should be removed or reused?
- How should Netflix balance relevance, diversity and novelty?
- How should recommendations change when the member's interests change?

Netflix's earlier engineering work describes homepage personalization as a multi-stage process involving candidate generation, filtering, ranking, row selection and page-level optimization.[\[netflixtechblog\]](#)

Therefore, the problem is better stated as:

Given a member, their previous behaviour, the available catalogue and the current context, what should Netflix display, where should it appear and in what order?

2.2 Why is the problem important?

A large catalogue is useful only when members can navigate it efficiently. If members repeatedly fail to find relevant content, they may become frustrated, reduce usage or cancel their subscription.

Netflix's 2015 peer-reviewed study reported that recommendations influenced approximately 80% of viewing hours at that time and estimated more than \$1 billion in annual value, mainly through reduced member churn. These are historical findings and should not be presented as current 2026 measurements.[[netflixtechblog](#)]

Recommendation contributes to:

- Faster content discovery.
- Better user experience.
- Higher engagement.
- Greater content utilization.
- More effective presentation of niche content.
- Improved retention.
- More efficient navigation of the catalogue.

However, clicks alone are not enough. A recommendation that produces a click but is abandoned after a few minutes may be less valuable than a recommendation that is completed and improves long-term satisfaction.

3. Data Analytics Solution

3.1 Overall recommendation architecture

Netflix does not use one single algorithm for every recommendation surface. Different surfaces require different objectives and models.

A simplified architecture consists of:

1. Data collection.
2. Data cleaning and feature construction.
3. Candidate generation.
4. Collaborative and content-based scoring.
5. Contextual and sequential ranking.
6. Diversity and availability filtering.
7. Homepage construction.
8. Online experimentation.




Figure 2. Multi-stage architecture of a Netflix-style personalization system.

Source: Amatriain, "Recommender Systems in Industry: A Netflix Case Study."

3.2 Evolution of Netflix recommendation

Table 2. Evolution of recommendation approaches

Approach	Main idea	Function
----------	-----------	----------

User-based Collaborative Filtering	Find members with similar behaviour	Recommend titles liked by similar members
Item-based Collaborative Filtering	Find relationships between titles	Recommend titles related to watched content
Matrix Factorization	Learn latent member and title vectors	Estimate member–title affinity
Specialized ML models	Train models for specific product surfaces	Support ranking and page construction
Recommendation Foundation Model	Centralize long-term preference learning	Provide shared embeddings and predictions
Generative recommendation	Predict recommendation sequences	Model sequential behaviour
GenRec	Use an LLM-backed ranking model	Rank catalogue items using rich context
GenPage	Generate homepage rows or entities	Construct a personalized page as a sequence

Netflix’s 2025 Foundation Model research explains that the company had accumulated numerous specialized models and sought to centralize preference learning to improve reuse and reduce maintenance costs.[\[netflixtechblog\]](#)

3.3 Data preparation

Before model training, interaction events need to be converted into reliable analytical data.

Important steps include:

- Removing duplicate or corrupted events.
- Distinguishing intentional actions from accidental clicks.
- Converting watch duration into engagement signals.
- Handling incomplete sessions.
- Applying regional catalogue restrictions.
- Separating profiles within an account.
- Using time-based training and testing splits.

- Restricting access to personal data.
- Applying retention and security policies.
- Creating features for recency and frequency.

Netflix largely relies on implicit feedback rather than only explicit ratings. A member may never give a numerical rating, but watching most of a film or repeatedly viewing a series can provide strong evidence of interest.

4. Collaborative Filtering

4.1 Concept

Collaborative Filtering recommends items using patterns of interaction among members and titles.

Its basic assumption is:

Members who have demonstrated similar preferences in the past may have similar preferences in the future.

For Netflix:

- Users represent member profiles.
- Items represent films and shows.
- Interactions represent watches, durations, searches, likes and other behaviour.

Table 3. Simplified user–item matrix

Member	Movie A	Movie B	Movie C	Movie D
User 1	5	4	?	1
User 2	5	?	4	1
User 3	?	2	5	4

The question marks represent unknown preferences. In a practical Netflix system, values would be derived from implicit signals rather than necessarily being explicit ratings.

4.2 User-based Collaborative Filtering

User-based Collaborative Filtering follows these steps:

1. Identify members with similar viewing patterns.
2. Find titles enjoyed by those members.
3. Recommend suitable titles to the target member.

For example, if two members repeatedly watch similar crime dramas and one of them watches a new investigative series, the system may recommend the series to the other.

Advantages

- Easy to understand.
- Naturally personalized.
- Useful for small and medium-sized datasets.
- Provides intuitive explanations based on similar members.

Limitations

- Pairwise member comparison is expensive at scale.
- Similarity is unreliable when members have few common interactions.
- Shared profiles can introduce noise.
- It does not naturally understand the order of viewing behaviour.
- Recommendations can change rapidly as new interactions arrive.

4.3 Item-based Collaborative Filtering

Item-based Collaborative Filtering focuses on title relationships.

For example:

Members who watched Movie A also frequently watched Movie B.

When a member watches Movie A, Movie B may be recommended.

Advantages

- Item relationships can be precomputed.
- Recommendations can be generated efficiently.
- Results are relatively stable.

- It supports explanations such as “Because you watched...”

Limitations

- Popularity bias can become stronger.
- Recent changes in member interest may be missed.
- Session context is weakly represented.
- Recommendations can become repetitive.
- Newly released titles have little interaction history.

4.4 Matrix Factorization

Matrix Factorization provides a scalable model-based method.

Let R be the user–item interaction matrix. The model approximates it as:

$$R \approx UV^T$$

where:

- U is the member latent-factor matrix.
- V is the title latent-factor matrix.
- Each row represents a lower-dimensional vector.
- The dot product estimates affinity.

For a member u and title i :

$$\hat{r}_{ui} = p_u^T q_i$$

where p_u is the vector for member u , q_i is the vector for title i , and $\hat{r}_{ui}$ is the estimated preference.

					
Member 1	✓		✓	✓	
Member 2		✓			✓
Member 3	✓	✓	✓		
Member 4				✓	✓

 $\approx$

	p_u	q_i	$\hat{r}_{ui}$	q_i	q_i
Member 1	1	.1	.88	-1.08	0.9
Member 2	-1	0	-0.9	1.0	-1.0
Member 3	.2	-1	0.38	0.6	1.2
Member 4	.1	1	-0.11	-0.9	-0.9

Figure 3. Matrix Factorization decomposes a sparse user–item matrix into latent member and title factors.

The latent dimensions may represent hidden preferences for:

- Action.
- Romance.
- Comedy.
- International content.
- Older films.
- Short or long-form storytelling.
- Specific narrative styles.

The dimensions are learned from interaction data and do not need to be manually named.

4.5 Netflix Prize

The Netflix Prize was conducted between 2006 and 2009. Netflix offered a \$1 million prize for improving its Cinematch algorithm by 10% on a held-out rating dataset.

The winning solution combined matrix factorization, neighbourhood methods and other techniques. The competition demonstrated the importance of model combination and strongly influenced recommender-system research.[\[netflixtechblog\]](#)

The Netflix Prize solution should not be described as Netflix’s current production algorithm. The competition used a historical rating dataset and a specific prediction objective. Netflix’s current system must also address:

- Implicit feedback.
- Real-time ranking.
- Regional availability.
- Sequential behaviour.
- Diversity.
- Cold start.
- Homepage construction.
- Long-term member value.
- Serving latency.

5. Suitability of Collaborative Filtering

5.1 Advantages

Table 4. Advantages of Collaborative Filtering

Advantage	Explanation
Personalization	Learns from individual member behaviour
Minimal manual feature definition	Behavioural relationships can be learned automatically
Sparse-data compatibility	Designed for incomplete user–item matrices
Hidden preference discovery	Latent factors capture complex taste patterns
Reusable relationships	Item relationships can be precomputed
Strong research foundation	Netflix Prize demonstrated the value of combined models
Relatively understandable	Similar members or titles can explain recommendations

5.2 Limitations

Cold start

A new member has little or no viewing history. A new title has few or no interactions. Collaborative Filtering therefore has insufficient evidence.

Netflix’s Foundation Model work addresses this by combining interaction data with metadata and content-based representations.[[netflixtechblog](#)]

Sparsity

Most members interact with only a small portion of the catalogue, leaving many user–item relationships unknown.

Popularity bias

A model trained primarily on historical interaction volume may repeatedly recommend already popular titles.

Weak sequential understanding

Classical Matrix Factorization does not naturally understand a sequence such as:

Watched a thriller yesterday

- searched for comedy today
- browsing on a mobile device
- currently wants a short title

Objective mismatch

Prediction accuracy is not identical to user satisfaction. Netflix must consider relevance, completion, diversity, novelty, retention and long-term value.

Privacy and responsible-use risks

Personalized recommendations process detailed behavioural information. Important risks include:

- Inferring sensitive interests.
- Collecting more information than necessary.
- Lack of transparency.
- Unintended bias.
- Excessive personalization.
- Inadequate control over recommendations.
- Retaining historical behaviour unnecessarily.

6. Modern Netflix AI Integration

6.1 Recommendation Foundation Model

Netflix's Foundation Model for Personalized Recommendation centralizes preference learning from:

- Comprehensive interaction histories.
- Content metadata.
- Contextual signals.

- Member and title representations.

The model is transformer-based and uses an autoregressive objective inspired by language modelling. It treats interactions as sequences rather than independent events. [[netflixtechblog](#)]

For example:

Watched A

- Watched B
- Searched C
- Watched D
- Liked D

The model can produce:

- Member embeddings.
- Title embeddings.
- Predictions.
- Features for downstream models.
- Representations for candidate generation.
- Fine-tuned models for specific applications.




Figure 4. Combining metadata-based and interaction-based title representations for recommendation cold start.

Source: Hsiao, Feng and Lamkhede, “Foundation Model for Personalized Recommendation,” Netflix Technology Blog, 2025.

Netflix also uses methods such as:

- Sparse attention.
- Sliding-window sampling.
- Incremental training.
- Metadata-based representations.
- Efficient inference.

These methods are necessary because active members may have thousands of interactions, while recommendation services often require millisecond-level responses.

6.2 Generative recommendation

Netflix has investigated models that generate recommendation sequences rather than only predicting independent title scores.

Netflix reported experiments scaling generative recommendation models from approximately 2 million to 1 billion backbone parameters. The experiments showed that larger models can improve reported recommendation tasks, but model size alone does not solve production problems. Training cost, decoding cost, latency, freshness and cold start remain important.[\[finance.yahoo\]](#)

This leads to an important engineering principle:

An algorithm must be evaluated by both quality and operational feasibility.

6.3 GenRec

GenRec is an LLM-backed recommendation ranker developed by Netflix.

It:

- Represents member history, item information and context as text.
- Uses a Netflix-adapted foundation LLM.
- Post-trains the model for ranking.
- Uses a catalogue-aware scoring head.

- Optimizes recommendation quality and longer-term member value.
- Uses efficient inference rather than generating unrestricted text at serving time.

Netflix reported approximately a 1.6% improvement in Mean Reciprocal Rank in one offline configuration while using approximately 40 times fewer Phase-2 labelled examples than the comparison setup. It also reported statistically significant improvements over a mature production ranker in a large online test.[[au.variety](#)]

These results are important, but they should not be generalized into the claim that every LLM-based recommender is better than every Collaborative Filtering model. The results come from Netflix’s specific model, data, infrastructure and evaluation design.



Figure 5. GenRec uses a foundation LLM, recommendation-specific post-training and catalogue-aware ranking.

Source: Li, Rao and Sehgal, “GenRec: Towards LLM-Native Recommendation at Netflix,” Netflix Technology Blog, 2026.

6.4 GenPage

Traditional homepage construction can be expressed as:

Candidate generation

- ranking
- re-ranking
- row selection
- page construction

GenPage treats the homepage as a structured sequence and generates rows or entities one at a time while considering member context and the content already placed on the page.[[netflixtechblog](#)]




Figure 6. GenPage generates a personalized homepage as a structured sequence of rows and entities.

Source: Wang, Pan and Baltrunas, “GenPage: Towards End-to-End Generative Homepage Construction at Netflix,” *Netflix Technology Blog*, 2026.

Netflix reported that one GenPage configuration produced an approximately 0.24% improvement in a core engagement metric and approximately 20% lower end-to-end serving latency than the mature production homepage recommender. These should be described as results from a reported online experiment rather than as universal improvements.

7. Critical Analysis

7.1 Is Collaborative Filtering still the best algorithm?

Collaborative Filtering remains an important foundation, but it is not sufficient as Netflix’s only algorithm.

Table 5. Comparison of recommendation algorithms

Algorithm	Main strength	Main limitation	Suitability for Netflix
User-based CF	Intuitive member similarity	Expensive pairwise computation	Low as a standalone system
Item-based CF	Reusable title relationships	Weak sequence understanding	Useful component
Matrix Factorization	Efficient latent representation	Cold start and weak context	Strong baseline

Content-based filtering	Helps new titles	May over-specialize	Strong complement
Sequence Transformer	Models order and recency	High computational cost	Very strong
Foundation Model	Shared representations	Expensive training and serving	Strong for large-scale use
GenRec	LLM-based contextual ranking	Infrastructure complexity	Promising
GenPage	Optimizes the complete page	Structured generation is complex	Promising for homepage construction

7.2 Scalability

User-based Collaborative Filtering

If every member is compared with every other member, naïve similarity computation can approach:

$$O(U^2)$$

where U is the number of members.

This is impractical at Netflix's scale without approximate nearest-neighbour search, clustering, sampling or offline computation.

Matrix Factorization

If $|\Omega|$ is the number of observed interactions and K is the number of latent factors, one SGD training pass is approximately proportional to:

$$O(|\Omega|K)$$

This is more practical because the model focuses on observed interactions rather than every possible user-item combination.

Transformer models

Self-attention traditionally has approximately:

$O(L^2)$

complexity with respect to sequence length L . Netflix therefore uses sparse attention, sliding-window sampling and efficient inference.

Table 6. Computational trade-offs

Approach	Main challenge	Possible solution
User-based CF	Pairwise member comparison	Approximate similarity search
Item-based CF	Large title relationship matrix	Offline precomputation and pruning
Matrix Factorization	Many observed interactions	Distributed training
Transformer	Long interaction histories	Sparse attention and sliding windows
LLM ranking	Large inference cost	Candidate restriction and caching
Generative homepage	Structured generation cost	Constrained decoding and efficient serving

7.3 Cold-start problem

New member

No viewing history

- weak collaborative representation
- uncertain recommendations

Potential solutions include:

- Asking members to choose initial genres or titles.
- Using regional and language preferences.
- Using broad but diverse starter recommendations.
- Updating the profile quickly from early interactions.
- Combining content and collaborative signals.

New title

No interaction history

→ weak collaborative evidence

Potential solutions include:

- Genre and thematic metadata.
- Cast and creator information.
- Language and country.
- Synopsis and trailer representations.
- Similarity to existing titles.
- Controlled exploration.

Netflix's Foundation Model research explicitly discusses metadata-based representations for unseen entities.[\[netflixtechblog\]](#)

7.4 Popularity, diversity and fairness

A system optimized only for immediate engagement can over-recommend popular titles. This may reduce exposure to niche, regional or less-established content.

Netflix's homepage research discusses the importance of balancing relevance with diversity, novelty, exploration and familiarity.[\[netflixtechblog\]](#)

A conceptual utility function can be written as:

$$\text{Utility} = \alpha(\text{relevance}) + \beta(\text{engagement}) + \gamma(\text{diversity}) + \delta(\text{novelty}) - \lambda(\text{repetition})$$

The values of the coefficients would vary by surface and business objective.

A responsible recommendation system should also monitor whether it systematically under-recommends certain languages, regions, creators or categories because of historical popularity differences.

7.5 Evaluation

Offline evaluation

Possible metrics include:

- Precision.
- Recall.
- Hit Rate.
- Mean Reciprocal Rank.
- NDCG.
- Coverage.
- Diversity.
- Novelty.

Mean Reciprocal Rank can be expressed as:

$$\text{MRR} = 1/N \sum_j 1/\text{rank}_j$$

where rank_j is the position of the first relevant title for member j .

Online evaluation

Netflix can use controlled experiments to measure:

- Playback starts.
- Completion.
- Session quality.
- Repeat visits.
- Search reduction.
- Long-term engagement.
- Retention.
- Satisfaction indicators.

A strong production evaluation should report:

1. Offline ranking performance.
2. Online member impact.
3. Long-term business impact.
4. Latency.
5. Infrastructure cost.
6. Freshness.
7. Diversity.
8. Cold-start performance.

A model that improves clicks but reduces completion or long-term satisfaction may not be a better model.

8. Proposed Hybrid Solution

8.1 Recommended architecture

The most suitable architecture for Netflix is a hybrid model combining:

- Collaborative signals.
- Content metadata.
- Sequential behaviour.
- Contextual features.
- Diversity constraints.
- Online experimentation.




Figure 7. The Netflix deep learning ranker for search personalization

Source: PyImageSearch, "Netflix Movies and Series Recommendation Systems," 2023.

Example:

Causal Inference provides formal tools to tease out the true **incremental** value of an **impression** for each profile: Heterogeneous Treatment Effect (**HTE**)





Figure 8. Causal Inference at Netflix

Source: PyImageSearch, "Netflix Movies and Series Recommendation Systems," 2023.

Table 7. Role of each component

Component	Purpose
Item-based CF	Generate titles related to recently watched content
Matrix Factorization	Learn broad member–title affinities
Content-based model	Handle new titles and metadata similarity
Sequence model	Understand order, recency and changing intent
Context model	Use device, locale, time and current session
Re-ranking layer	Balance relevance, diversity and novelty

Page-generation model	Construct coherent homepage rows
Experimentation layer	Measure actual member and business impact

8.2 Example

Consider a member who:

1. Recently watched two crime dramas.
2. Abandoned a long historical series.
3. Searched for a comedy.
4. Uses Netflix on a mobile device in the evening.

A Collaborative Filtering model may identify titles watched by members with similar histories. A content-based model may find titles with similar crime-drama metadata. A sequence model may identify that the member’s recent intent has shifted toward comedy. The final ranking layer may choose a short comedy that is regionally available, avoid repeating similar titles and place it in an appropriate row.

This example demonstrates why Netflix requires more than a static user–item matrix.

9. Overall Assessment

Table 8. Technical and business assessment

Criterion	Assessment
Personalization	Very strong
Relevance	Strong, especially with sequence and content signals
Cold start	Improved through metadata and content representations
Scalability	Matrix Factorization is efficient; foundation models require specialized infrastructure
Explainability	Classical CF is easier to explain than LLM-based generation
Diversity	Requires explicit objectives and monitoring

Freshness	Requires incremental updates and recent-session features
Latency	Candidate generation, caching and efficient serving are essential
Maintainability	Shared models can reduce duplicated systems but increase infrastructure complexity
Business value	Potential improvement in discovery, engagement and retention

10. Responsible Use of AI

Perplexity was used in preparing this report to assist with:

- Finding Netflix technical publications.
- Organizing the literature survey.
- Comparing recommendation algorithms.
- Improving the structure and readability.
- Identifying suitable visual references.
- Separating historical and current claims.

The information was verified against primary or authoritative sources, including:

- Netflix Research.
- Netflix Technology Blog.
- Netflix Investor Relations.
- The U.S. Securities and Exchange Commission.
- ACM and IEEE publications.

The following verification decisions were made:

- Historical statistics were labelled as historical.
- The Netflix Prize winner was not described as Netflix’s current algorithm.
- Unsupported data-volume estimates were excluded.
- Experimental results from GenRec and GenPage were not generalized beyond their reported evaluations.
- The official SEC filing was used for the 2025 revenue figure.
- Visuals were selected from official or academically credible sources where possible.
- Conceptual third-party diagrams were explicitly labelled as conceptual rather than official Netflix architecture.

11. Conclusion

Netflix is an excellent industry application for studying Data Analytics and Collaborative Filtering.

Collaborative Filtering models relationships between members and titles. User-based filtering uses similar members, item-based filtering uses similar titles, and Matrix Factorization learns latent representations. These methods are highly relevant to the historical and conceptual development of recommender systems.

However, Collaborative Filtering alone cannot solve all of Netflix's current problems. It struggles with cold start, sparsity, popularity bias, changing interests, limited sequence understanding and the difference between prediction accuracy and long-term member satisfaction.

Netflix's recent research shows a movement toward transformer-based recommendation foundation models, generative recommendation, GenRec and GenPage. These systems can incorporate long-term histories, recent activity, content metadata, context and catalogue-level constraints.

The most accurate conclusion is therefore:

Collaborative Filtering remains an important foundation for Netflix's recommendation problem, but a modern Netflix-style system should combine Collaborative Filtering with content-based modelling, sequential deep learning, contextual ranking, diversity controls and experimentation.

The main industrial lesson is that recommendation quality is not defined only by prediction accuracy. A successful system must also be scalable, responsive, fresh, diverse, responsible and aligned with long-term member value.

References

[1] Netflix, "How Netflix's Recommendations System Works," Netflix Help Center, 2026. Available: <https://help.netflix.com/en/node/100639>.

[2] Netflix, Inc., "Annual Report on Form 10-K for the Fiscal Year Ended December 31, 2025," U.S. Securities and Exchange Commission, 2026.

- [3] Netflix, Inc., “2026 Proxy Statement,” Netflix Investor Relations, 2026.
- [4] K.-J. Hsiao, Y. Feng, and S. Lamkhede, “Foundation Model for Personalized Recommendation,” Netflix Technology Blog, 2025. Available: <https://netflixtechblog.com/foundation-model-for-personalized-recommendation-1a0bd8e02d39>.
- [5] M. Bhattacharya and S. D. Lamkhede, “Augmenting Netflix Search with In-Session Adapted Recommendations,” in *Proceedings of the 16th ACM Conference on Recommender Systems (RecSys '22)*, 2022, doi: 10.1145/3523227.3547407.
- [6] O. Khoo, “Picturing Diversity: Netflix’s Inclusion Strategy and the Netflix Recommender Algorithm (NRA),” *Television & New Media*, vol. 24, no. 3, pp. 281–297, 2023, doi: 10.1177/15274764221102864.
- [7] H. Wang, “MovieMat: Context-aware Movie Recommendation with Matrix Factorization by Matrix Fitting,” arXiv preprint arXiv:2204.13003, 2022.